

Prediction of ozone level in Boston

*Lucas Emanuel Resck Domingues**


Lucas Machado Moschen†

TO DO

1. Week model: the same analysis
2. Escrever melhor o texto em algumas transições.
3. Alguma combinação de modelo?
4. Residual analysis with adjusted error, not predictive one.

Load and visualize

Daily average level of O3 in Boston



Data treatment

We noticed that some days do not exist in the dataset, for example, the day August 31, 2001 does not have information in the dataset.

##	X	City	State	Site.Num	Date.Local	O3.Mean
## 148	148	Boston	Massachusetts	42	2001-08-28	0.024583
## 149	149	Boston	Massachusetts	42	2001-08-29	0.015000
## 150	150	Boston	Massachusetts	42	2001-08-30	0.022333
## 151	151	Boston	Massachusetts	42	2001-09-01	0.021958

*Escola de Matemática Aplicada

†Escola de Matemática Aplicada

```
## 152 152 Boston Massachusetts      42 2001-09-02 0.018750
## 153 153 Boston Massachusetts      42 2001-09-03 0.028708
```

Also, there is duplicated days, as June 9, 2002:

```
##      X   City      State Site.Num Date.Local  O3.Mean
## 412 412 Boston Massachusetts      42 2002-06-08 0.022917
## 413 413 Boston Massachusetts      42 2002-06-09 0.036190
## 414 414 Boston Massachusetts      42 2002-06-09 0.037000
## 415 415 Boston Massachusetts      42 2002-06-10 0.023389
```

The duplicated one is easier to deal, but the missing values are harder. First we calculate the mean value between the duplicated.


The rate of missing values is almost 5% of the dataset.

```
## [1] 0.04453367
```

So as to solve that problem, we make a knn imputation using the month ($k = 30$)

```
o3.clean <- knn.impute(as.matrix(o3.ts), k = 30)
o3.clean <- as.ts(o3.clean)
plot(o3.clean, main = 'Daily average level of O3 in Boston (after imputation)',
      xlab = 'Date', ylab = 'Kg/m3')
```

Daily average level of O3 in Boston (after imputation)



We also separate our data between training and test because of modelling best practices.

```
o3_train = o3.clean[1:(length(o3)[1] - 365),]
o3_test = o3.clean[(length(o3)[1] - 365 + 1):length(o3)[1],]
```

Models: case 1

Now we develop some models using the train data.

The metric to compare is the Mean Absolute Error (MAE) in the predictions:

```
mae <- function(ytrue, ypred)
{
  return(mean(abs(ytrue - ypred)))
}
```


We will use `rollapply` in order to calculate the error, considering the last two years to predict one week forward.

Baseline Model

We will do the naive forecast to the baseline model.

```
## [1] 0.008015428
```

Baseline model prediction



Decompose

First of all we make a seasonality test using Kruskal-Wallis. Actually it tests whether samples originate from the same distribution. We can organize it to be samples for each corresponding day. We compare two different frequencies: monthly and yearly. The second one showed the smallest p-value, in particular less than 0.05. For that reason, we will use 365 as the seasonality.

```
##
## Kruskal-Wallis rank sum test
##
## data:  o3_train and g
## Kruskal-Wallis chi-squared = 32.983, df = 30, p-value = 0.3233
##
## Kruskal-Wallis rank sum test
##
## data:  o3_train and g
```


```
## Kruskal-Wallis chi-squared = 2122.9, df = 364, p-value < 2.2e-16
```

Additive model


First we analyse the MAE.

```
## [1] 0.007963981
```

Additive decompose prediction




We also can fit the model using `t.window` and analyse the reminder of the method.




The ACF and the PACF of the remainder:

ACF Remainder



PACF Remainder




We see that there are a big spike when lag = 365. It seems not so good for a remainder. We could fit an ARMA model in this remainder yet.

Multiplicative model


First we analyse the MAE.

```
## [1] 0.00795386
```

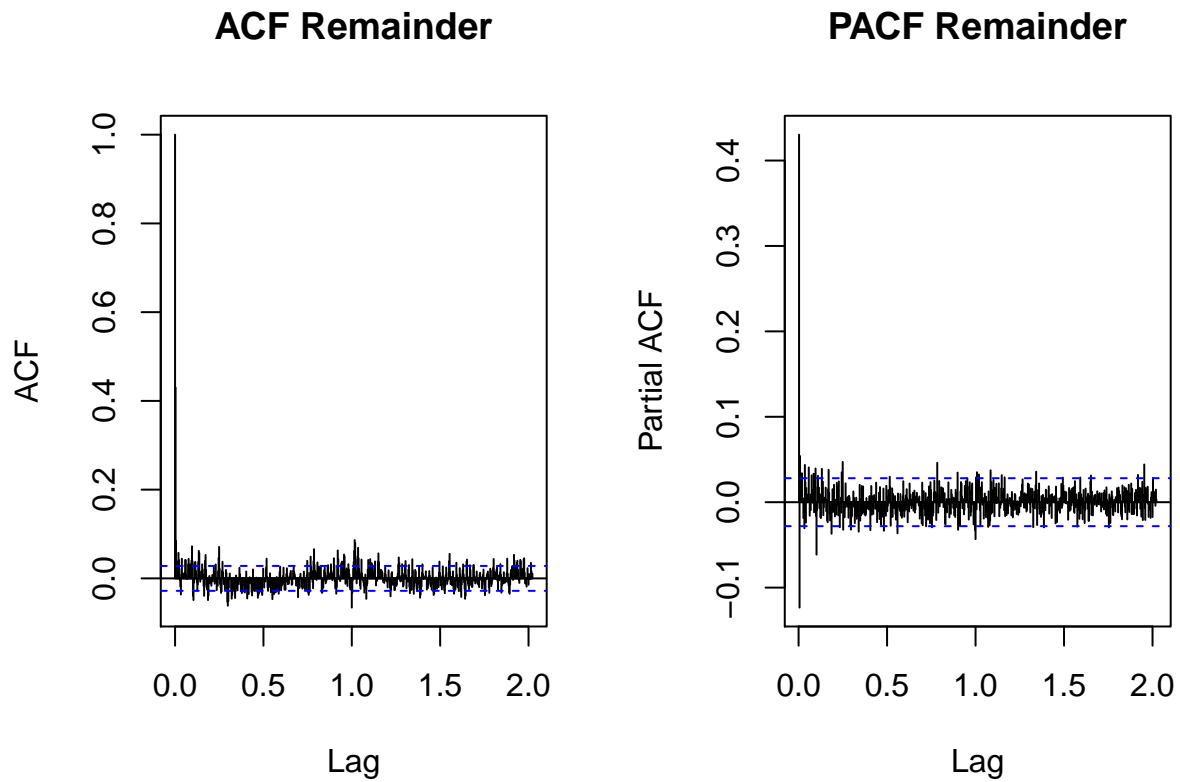
Multiplicative decompose prediction



We also can fit the model using `t.window` and analyse the reminder of the method.



The ACF and the PACF of the reminder:




We see that there are a big spike when lag = 365. It seems not so good for a reminder. We could fit an ARMA model in this reminder yet. The same problem as before.

Regression

We tested for seasonalities, and we settled with 365. So we will fit a regression model with the seasonality dummies. We see the MAE:

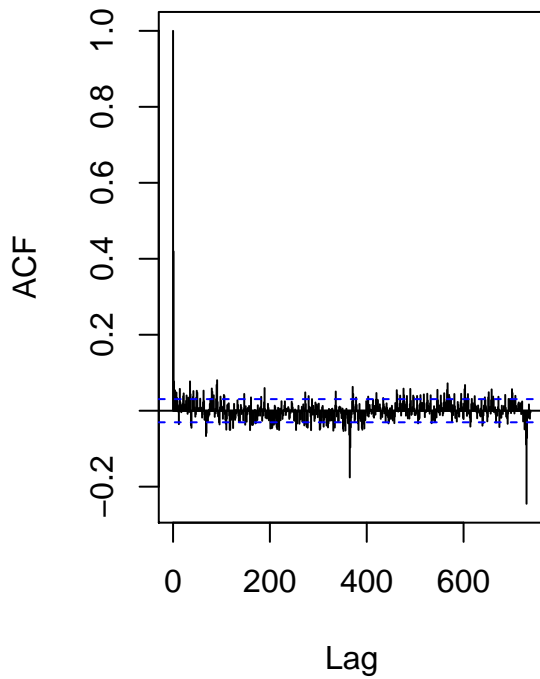
```
## [1] 0.007518997
```

Regression prediction

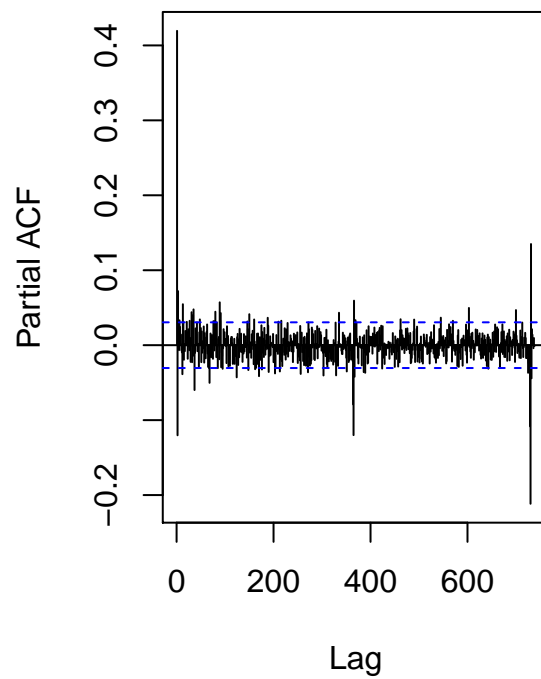


Now we analyse the residuals. The analysed residuals will be those from each model in the rolling window. It's valid because every model has (the assumption of) white noise with the same variance. Let's see the ACF and PACF:

ACF Remainder



PACF Remainder



As before, we see spikes in lag = 365, 730. We expect a WN to not have this.


Holt-Winters

Now we will try Holt-Winters models. In fact, because of apparently seasonality, we will consider complete Holt-Winters models, both additive and multiplicative.


Additive

```
## [1] 0.03493751
```

Additive Holt-Winters prediction



MAE is not so good, nor is the resulting graph. HW tends to “explode”. We have seen better results. Let’s analyse the residuals:




We see the same problems as before: high correlated lag = 365, evidence of this not being a WN.

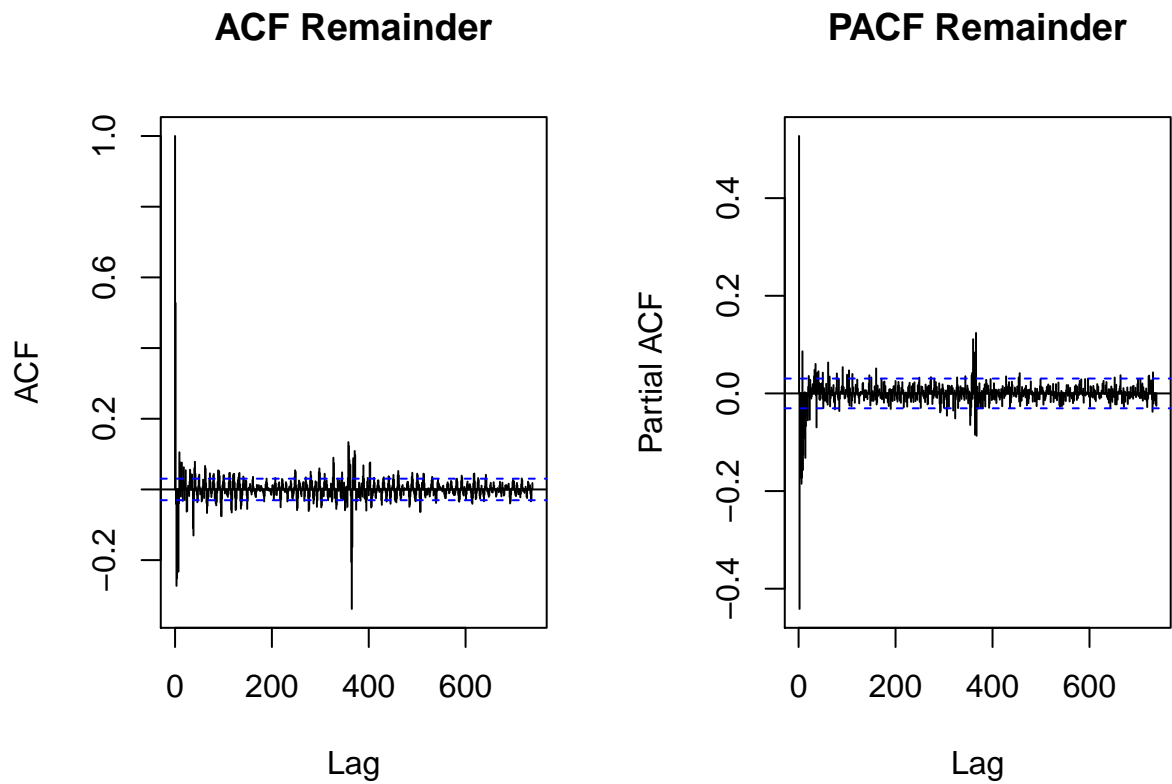
Multiplicative

```
## [1] 0.0348543
```

Multiplicative Holt–Winters prediction




We see also a not so good MAE and graph. Let's analyse the residuals:



Same problems as before. Holt-Winters doesn't appear to work well in this case.

ARMA

We can see the ACF and PACF:



Based on these graphs, we see both graphs has a exponentially decay, the first after the $p - q = 1$ or $p - q = 2$. In order to identify the model, we will compare the adjusted ARMA models with different p and q . First we simply fit it to look at the Akaike Information Criteria (AIC) and the significance of the parameters estimated.

The AIC measures the goodness of fit and the simplicity of the model into a single statistic. Generally we aim to reduce the AIC.

$$AIC = 2k - 2\ln(\hat{L}),$$

where $k = p + q + 2$ and \hat{L} is the maximum value of the likelihood for the model.

```
##
## Call:
## arma(x = o3_train, order = c(2, 1))
##
## Model:
## ARMA(2,1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0256895 -0.0050812 -0.0003623  0.0043066  0.0398679
##
## Coefficient(s):
##              Estimate Std. Error t value Pr(>|t|)
## ar1          1.4018803   0.0150148   93.367 < 2e-16 ***
## ar2         -0.4073756   0.0146509  -27.805 < 2e-16 ***
## ma1         -0.9294688   0.0059259 -156.849 < 2e-16 ***
## intercept    0.0001116   0.0000292    3.822 0.000132 ***
## ---
```

```

## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Fit:
## sigma^2 estimated as 5.396e-05,  Conditional Sum-of-Squares = 0.26,  AIC = -34030.29
##
## Call:
## arma(x = o3_train, order = c(3, 1))
##
## Model:
## ARMA(3,1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0280340 -0.0049450 -0.0003378  0.0043396  0.0391002
##
## Coefficient(s):
##      Estimate Std. Error t value Pr(>|t|)
## ar1      1.436e+00  1.616e-02  88.837 < 2e-16 ***
## ar2     -5.972e-01  2.357e-02 -25.337 < 2e-16 ***
## ar3      1.537e-01  1.485e-02  10.346 < 2e-16 ***
## ma1     -9.060e-01  8.615e-03 -105.163 < 2e-16 ***
## intercept 1.531e-04  3.928e-05   3.897 9.75e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Fit:
## sigma^2 estimated as 5.282e-05,  Conditional Sum-of-Squares = 0.26,  AIC = -34132.4
##
## Call:
## arma(x = o3_train, order = c(3, 2))
##
## Model:
## ARMA(3,2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.027737 -0.004941 -0.000356  0.004364  0.039341
##
## Coefficient(s):
##      Estimate Std. Error t value Pr(>|t|)
## ar1      1.2360037  0.0891579  13.863 < 2e-16 ***
## ar2     -0.3130929  0.1287091  -2.433 0.014992 *
## ar3      0.0686044  0.0422242   1.625 0.104213
## ma1     -0.7035439  0.0887237  -7.930 2.22e-15 ***
## ma2     -0.1896764  0.0832069  -2.280 0.022633 *
## intercept 0.0001719  0.0000448   3.837 0.000124 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Fit:
## sigma^2 estimated as 5.279e-05,  Conditional Sum-of-Squares = 0.26,  AIC = -34133.17
##

```

```

## Call:
## arma(x = o3_train, order = c(4, 3))
##
## Model:
## ARMA(4,3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0275745 -0.0049606 -0.0003479  0.0043427  0.0391448
##
## Coefficient(s):
##      Estimate Std. Error t value Pr(>|t|)
## ar1      0.9060239   0.3974686   2.279  0.02264 *
## ar2     -0.0265556   0.4953457  -0.054  0.95725
## ar3      0.1412857   0.1831460   0.771  0.44045
## ar4     -0.0326660   0.0554639  -0.589  0.55589
## ma1     -0.3723047   0.3971375  -0.937  0.34852
## ma2     -0.2997470   0.2906608  -1.031  0.30242
## ma3     -0.1781035   0.1224956  -1.454  0.14596
## intercept  0.0002418   0.0000921   2.625  0.00866 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Fit:
## sigma^2 estimated as 5.281e-05, Conditional Sum-of-Squares = 0.26, AIC = -34127.1
##
## Call:
## arma(x = o3_train, order = c(4, 2))
##
## Model:
## ARMA(4,2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0278979 -0.0049234 -0.0003225  0.0043674  0.0393150
##
## Coefficient(s):
##      Estimate Std. Error t value Pr(>|t|)
## ar1      1.059e+00   4.642e-01   2.281  0.0226 *
## ar2     -6.008e-02   6.680e-01  -0.090  0.9283
## ar3     -4.541e-02   2.895e-01  -0.157  0.8754
## ar4      3.665e-02   8.276e-02   0.443  0.6579
## ma1     -5.259e-01   4.625e-01  -1.137  0.2555
## ma2     -3.481e-01   4.171e-01  -0.834  0.4040
## intercept  2.042e-04   9.095e-05   2.246  0.0247 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Fit:
## sigma^2 estimated as 5.281e-05, Conditional Sum-of-Squares = 0.26, AIC = -34128.87


```

First we see that in the last two model, there is no statistical significance in the major part of the parameters, so we'll no consider it. The first, second and third models have significant parameters and similar AIC. In especial the 2° and 3° has the smallest AIC. So we will compare them both.

We see the MAE of the ARIMA(3, 0, 1):

```
## [1] 0.006463936
```

ARIMA(3,0,1) prediction




And the MAE of the ARIMA(3, 0, 2):

```
## [1] "MAE ARIMA(3,0,2)"
```


```
## [1] 0.006483882
```

ARIMA(3,0,2) prediction



The first model seems a little better. So, let's check the residuals to observe its problems.

```
model <- arima(o3_train, order = c(3,0,1))  
checkresiduals(model, lag = 2*freq, lag.max = 2*freq)
```

```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(3,0,1) with non-zero mean
## Q* = 1627, df = 725, p-value < 2.2e-16
##
## Model df: 5.    Total lags used: 730
```


It's interesting to note the ACF has a peak around the 365, so the ARMA model did not seem to capture the seasonality. It would be better to fit an Seasonal ARIMA further. The histogram is pretty similar to normal distribution. The test made analyses if the mean is 0. It may be because the p-value is really small. That's the reason to adapt ARMA with STL, that is, fit an ARMA model in the residuals.

Adapting ARMA

The ARMA seems to fit well as we can see so far. However, it's not capturing other characteristics on the data, as seasonality. For that reason, we will combine the stl and arma model and extract the best of each one. We will decompose the series in trend and seasonality and in the reminder, we fit an arima model with `auto.arima()`.

```
## [1] 0.007854377
```

STL + ARIMA prediction




However the MAE isn't improved by this model. For that reason, this model was disregarded.

Comparing models in the test data.

We chose some of the best models to compare with the testing data.


```
## [1] 0.006599825
```

Multiplicative decompose prediction




[1] 0.006393855

Regression prediction



[1] 0.0256431


Multiplicative Holt–Winters prediction




```
## [1] "MAE ARIMA(3,0,1)"
```

```
## [1] 0.004755896
```

ARIMA(3,0,1) prediction



Finally, we have this bar graphic to compare the values:




Models: case 2

In case two, we have to aggregate the diary days in a week, starting from the sunday, as requested. So we calculate the mean value in the week to be its representant. The models may be very similar to the previous. We may see less outliers. We also will separate train and test data. The first day in the data is April 1, 2001, a Sunday. SO we do not worry about that.

```
o3_week <- c(1:floor(length(o3.clean)/7))
for(i in seq(1,length(o3.clean)-7, 7)){
  o3_week[ceiling(i/7)] = mean(o3.clean[i:(i+6)])
}

plot(ts(o3_week), main = 'Weekly average level of O3 in Boston (after imputation)',
      xlab = 'Week', ylab = 'Kg/m3')
```

Weekly average level of O3 in Boston (after imputation)



```
o3_train_week = o3_week[1:(length(o3_week)[1] - 52)]  
o3_test_week = o3_week[-c(1:(length(o3_week)[1] - 52))]
```